

# Lab 3: Regression Competition

In the Hot New Jersey Night

*Paul Nguyen*

## **group\_F\_process**

We eliminate undesired variables and also automatically generate a new dataframe with all the variables squared and cubic(ed).

## **group\_F\_automatic**

We run regsubsets on the dataframe and generate a lm object automatically, no need of manually checking and writing the LM.

```
library(tidyverse)
library(leaps)

d <- read.csv("http://andrewpbray.github.io/data/crime-train.csv")

# Data wrangling
group_F_process <- function(training_data) {
  dw<-as_tibble(training_data)
  dw<-select (training_data,-c(state,county,community,communityname,LemasSwornFT,LemasSwFTPerPop,LemasS
  vars<-c()
  for (i in 1:(length(dw)-1)) {
    vars<- c(vars, names(dw)[i])
  }
  sqrd<-data.frame(lapply(vars, function(x){dw[,x]^(1/2)}))
  cubc<-data.frame(lapply(vars, function(x){dw[,x]^(1/3)}))
  names(sqrd)<-paste0(vars, "Sq")
  names(cubc)<-paste0(vars, "Cub")
  dw<-cbind(dw, sqrd,cubc)
  return(dw)
}

# Manually fits model
group_F_fit <- function(training_data) {
  m1 <- lm(ViolentCrimesPerPop~
    MalePctDivorceCub+
    PctKids2ParCub+
    PctIlleg+
    PctPersDenseHousSq+
    RentLowQ+
    MedRent,
    training_data)
  m1
}

# Gets MSE
group_F_MSE <- function(model, data) {
  mean((data$ViolentCrimesPerPop - predict.lm(model, data))^ 2)
```

```

}

# Automatically fits model
group_F_automated_fit <- function(data, met) {
  leaps<-regsubsets(ViolentCrimesPerPop~.,
                    data = data,
                    nvmax = 70,
                    method = met)
  best<-summary(leaps)$which[which.max(summary(leaps)$adjr2),]
  variables <- c()
  for (i in 2:length(best)) {
    if (best[i] == TRUE) {
      variables <- c(variables, names(best)[i])
    }
  }
  vars<- paste(variables, collapse = "+")
  formula<- paste("lm(ViolentCrimesPerPop ~ ",vars," , data =", deparse(substitute(data)),")")
  m1<-eval(parse(text=formula))
  return(m1)
}

dw <- group_F_process(d)
bestF <- group_F_automated_fit(dw, "forward")

## Warning in leaps.setup(x, y, wt = wt, nbest = nbest, nvmax = nvmax,
## force.in = force.in, : 1 linear dependencies found

## Reordering variables and trying again:
bestB <- group_F_automated_fit(dw, "backward")

## Warning in leaps.setup(x, y, wt = wt, nbest = nbest, nvmax = nvmax,
## force.in = force.in, : 1 linear dependencies found

## Reordering variables and trying again:
MSE_F <- group_F_MSE(bestF, dw)
MSE_B <- group_F_MSE(bestB, dw)

```

looks like backward is the better way to go

## Problem Set

2.

- a. The lasso, relative to least squares, is less flexible so that it reduces bias when there is collinearity in the data. It will give improved prediction accuracy when its increase in bias is less than its decrease in variance
- b. Similarly to the lasso, the ridge regression is less flexible than the least squares method so that it provides more stability to the model. It will give improved prediction accuracy when its increase in bias is less than its decrease in variance.
3. Suppose we estimate the regression coefficients in a linear regression model by minimizing  $\sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2$  subject to  $\sum_{j=1}^p |\beta_j| \leq s$

- (a) Note that increasing  $s$  means that our coefficients become larger and larger... they are less constrained. This means that the model will fit to the training data increasingly well. thus, as we increase  $s$  from 0, the training RSS will steadily decrease.
  - (b) I think that a similar reasoning can be applied to Test RSS. At  $s = 0$ , the model is just  $B_0 = \text{mean of the data}$ . As we increase  $s$ , the model becomes more and more complex, until  $s$  becomes large enough that it encompasses  $\hat{\beta}$  where  $\hat{\beta}$  is the least sum of squares model. If our test data is similar to our training data, test RSS should decrease as we increase  $s$ . However, I believe that the model may become too complex, and so Test RSS may eventually start decreasing when we start overfitting our data.
  - (c) For variance, we can think about this as model complexity. When  $s = 0$ , there is no variance. The beta terms are just 0. As we increase  $s$  (this is equivalent to decreasing  $\lambda$  in the lasso function), the beta terms start emerging, and variance should start to steadily increase.
  - (d) For squared bias, when  $s = 0$ , bias here is super big. Then when we increase  $s$ , we start to add model complexity and so bias just starts to decrease.
  - (e) I think here irreducible error remains constant.  $s$  really only affects the beta terms, which estimate  $y$ . However, irreducible error comes from the data not being perfect and having a little bit of jitter no matter what level we set  $s$  to.
4. Suppose we estimate the regression coefficients in a linear regression model by minimizing  $\sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2 + \lambda \sum_{j=1}^p \beta_j^2$
- (a) Here, this is a ridge regression whose complexity starts at the least squares method and then decreases as  $\lambda$  increases. Thus, when we increase  $\lambda$  from 0, the training RSS will decrease pretty steadily. The training RSS would have been the highest at  $\lambda = 0$ .
  - (b) When we repeat for test RSS, I believe that the test RSS will decrease and then eventually start increasing. My reasoning: At  $\lambda = 0$ , this is equivalent to the least squares solution.. it may be needlessly complex. When we increase  $\lambda$ , we start punishing for needless predictors, which will bring our test RSS down. However, when  $\lambda$  becomes too high, then having any predictors at all will punish our model, and then our bias is overpowering any decreases in variance that we brought.
  - (c) Increasing  $\lambda$  decreases our model complexity. Thus, our variance will steadily decrease.
  - (d) Squared Bias will steadily increase as we increase  $\lambda$ . We decrease model complexity in exchange for introducing bias to our model.
  - (e) Again, irreducible error will remain constant. It is not affected by our model predictors.
5. given:  $n = 2, p = 2, x_{11} = x_{12}$  suppose that  $y_1 + y_2 = 0, x_{12} + x_{21} = 0$  so that the estimate for the intercept for a ridge regression, lasso, or least squares is 0.  $\hat{\beta}_0 = 0$
- (a) for ridge regression, minimize  $\sum_{j=1}^p (y_i - \beta_j)^2 + \lambda \sum_{j=1}^p \beta_j^2$   $((y_1 - \beta_1)^2 + (y_2 - \beta_2)^2) + ((\lambda\beta_1)^2 + (\lambda\beta_2)^2)$
  - (b) in this setting, ridge regression estimates take the form:  $\hat{\beta}_j^R = y_j/(1 + \lambda)$  (note talk to andrew about this)
  - (c) for lasso, minimize  $\sum_{j=1}^p (y_i - \beta_j)^2 + \lambda \sum_{j=1}^p |\beta_j|$   $((y_1 - \beta_1)^2 + (y_2 - \beta_2)^2) + (\lambda|\beta_1| + \lambda|\beta_2|)$
  - (d) lasso estimates:

$$\hat{\beta}_2 = y_2 - \lambda/2 \implies \hat{\beta}_2 = -y_1 - \lambda/2$$

$$\hat{\beta}_1 \neq \hat{\beta}_2$$

$$y_1 > \lambda/2 \implies y_2 < \lambda/2$$

$$\hat{\beta}_1 = y_1 - \lambda/2$$

$$\hat{\beta}_2 = y_2 + \lambda/2 \implies \hat{\beta}_2 = -y_1 + \lambda/2$$

Again,

$$\hat{\beta}_1 \neq \hat{\beta}_2$$

6.

- (a) For some  $y_1$  and  $\lambda > 0$ , plot 6.12 as a function of  $\beta_1$  Let  $y_1 = 2$  and  $\lambda = 1$

```

y1 <- 2
lambda <- -1
B1 <- seq(from = -10, to = 10, by = .01)
functionB1r <- (y1)^2 - 2*y1*B1 + B1^2 + lambda*B1^2
functionB1r2 <- (y1-B1)^2 + lambda*(B1^2)
min(functionB1r)

```

```
## [1] 2
```


```
min(functionB1r2)
```

```
## [1] 2
```

```

df <- data.frame(B1 = B1, functionB1r = functionB1r)
ggplot(data = df, mapping = aes(x = B1, y = functionB1r)) +
  geom_line() +
  geom_line(mapping = aes(x = y1/(lambda+1)))

```



note:  $\sum_{j=1}^p (y_i - \beta_j)^2 + \lambda \sum_{j=1}^p \beta_j^2$  represented by line, and line represents  $\beta_1^r = y_1/(1 + \lambda)$ .  $B_1$  is at min of function

(b) plot 6.13 as a function of  $\beta_1$

```

B1 <- seq(from = -3, to = 6, by = .01)
functionB1l <- (y1)^2 - 2*y1*B1 + B1^2 + abs(B1)
bayarea <- (y1 - B1)^2 + lambda*abs(B1)
min(bayarea)


```

```
## [1] 1.75
```


```
min(functionB1l)
```

```
## [1] 1.75
```

```
dfl <- data.frame(B1 = B1, functionB1l = functionB1l)
ggplot(data = dfl, mapping = aes(x = B1, y = functionB1l)) +
  geom_line(col = "purple2") +
  geom_line(mapping = aes(x = (y1 - (lambda/2))))
```



```
dfltest <- data.frame(B1 = B1, bayarea = bayarea)
ggplot(data = dfltest, mapping = aes(x = B1, y = bayarea)) +
  geom_line(col = "purple2") +
  geom_line(mapping = aes(x = (y1 - (lambda/2))))
```



## Additional Exercise

Using the `glmnet` package, construct a ridge regression and LASSO model to predict violent crime in the training data set. A description for how to use this package can be found in your book on pages 251 - 255.

How many variables were selected by the LASSO?

What are the training MSEs for ridge and LASSO using the optimal value of  $(\lambda)$ ?

If the MSEs differed, why do you think one is higher than the other in this setting?


```
library(glmnet)
```

```
## Loading required package: Matrix
##
## Attaching package: 'Matrix'
## The following object is masked from 'package:tidyr':
##
##     expand
## Loading required package: foreach
##
## Attaching package: 'foreach'
## The following objects are masked from 'package:purrr':
##
##     accumulate, when
```

```
## Loaded glmnet 2.0-18
grid <- 10^seq(10,-2, length =100)
xx <- model.matrix(ViolentCrimesPerPop ~., dw)[-1]
yy = dw$ViolentCrimesPerPop
ridge.mod = glmnet(x = xx, y = yy, alpha = 0,
                   lambda = grid)
dim(coef(ridge.mod))
```

```
## [1] 289 100
```

```
set.seed(11)
train <- sample(1:nrow(xx), nrow(xx)/2)
test <- (-train)
yy.test = yy[test]
cvouttr <- cv.glmnet(xx[train,], yy[train], alpha =0)
plot(cvouttr)
```



```
bestlamr <- cvouttr$lambda.min
bestlamr
```

```
## [1] 0.3237577
```


```
ridgepred = predict(ridge.mod, s = bestlamr, newx = xx[test,])
testmserridge <- mean((ridgepred - yy.test)^2)
testmserridge
```

```
## [1] 0.01849113
```

```
#now for lasso
lasso.mod = glmnet(x = xx, y = yy, alpha = 1,
                   lambda = grid)
dim(coef(lasso.mod))
```

```
## [1] 289 100
```

```
set.seed(11)
cvoutlasso <- cv.glmnet(xx[train,], yy[train], alpha =1)
plot(cvoutlasso)
```



```
bestlamlasso <- cvoutlasso$lambda.min
bestlamlasso
```

```
## [1] 0.005276443
```

```
lassopred = predict(lasso.mod, s = bestlamr, newx = xx[test,])
testmselasso <- mean((lassopred - yy.test)^2)
testmselasso
```

```
## [1] 0.05865972
```

```
lassocoeff = predict(lasso.mod, type = "coefficients", s = bestlamlasso)
lassocoeff[lassocoeff != 0]
```

```
## <sparse>[ <logic> ] : .M.sub.i.logical() maybe inefficient
```

```
## [1] 0.3691481182 -0.0570796828 -0.0007791917 0.0326701084 0.1450908366
## [6] -0.0077677206 0.1808884867 0.1246412015 -0.0159766045 0.0419122868
## [11] 0.0050247054 -0.0602653933 -0.3124524337 0.0025875656 0.0764096092
## [16] 0.0420028221 0.0028152885 0.1453804569
```

18 variables for lasso model. mse for ridge:.01849; mse for lasso: .0586 My guess for why the test MSE's differed was because the best lambda for lasso ended up being super small when compared to the lambda for ridge. This means that my lasso model ended up being more complex than the ridge model. Also, the lasso model should have eliminated any needless predictors, so since it was more complex, tried modeling the data more closely compared to the ridge model. However, it may have still been overfitting the test data, which could explain the higher test mse.



```

group_F_automated_fit <- function(training_data){
  mb <- lm(ViolentCrimesPerPop ~ population +
    racePctWhite +
    agePct12t29 +
    numbUrban +
    medIncome +
    PctEmploy +
    pctWSocSec +
    medFamInc +
    whitePerCap +
    PctEmploy +
    MalePctDivorce +
    MalePctNevMarr +
    PctWorkMom +
    PctIlleg +
    PctImmigRec5 +
    PctImmigRec8 +
    PersPerOccupHous +
    PctHousLess3BR +
    PctVacantBoarded +
    PctHousOccup +
    PctVacMore6Mos +
    RentLowQ +
    MedRent +
    MedOwnCostPctIncNoMtg +
    NumStreet,
    training_data)
  mb
}

```

```

group_F_process <- function(data){
  data %>%
    mutate(T_MPD = MalePctDivorce^(1/3),
           T_PK2P = PctKids2Par^(1/3),
           T_PPDH = PctPersDenseHous^2)
}

```

```

group_F_fit2.0 <- function(training_data) {
  m1 <- lm(ViolentCrimesPerPop ~ T_MPD + T_PK2P +
    PctIlleg + T_PPDH + RentLowQ+MedRent, training_data)
  m1
}

```

```

group_F_fit <- function(training_data) {
  m1 <- lm(ViolentCrimesPerPop~ I(MalePctDivorce^(1/3)) + I(PctKids2Par^(1/3)) +
    PctIlleg + (PctPersDenseHous^2) + RentLowQ+MedRent, training_data)
  m1
}

```

```

group_F_MSE <- function(model, data) {
  MSE<-mean((data$ViolentCrimesPerPop - predict.lm(model, data)) ^ 2)
  return (MSE)
}

```

```

library(leaps)
#forward
regsubsets(ViolentCrimesPerPop ~ population +
            householdsiz +
            racepctblack +
            racePctWhite +
            racePctAsian +
            racePctHisp +
            agePct12t21 +
            agePct12t29 +
            agePct16t24 +
            agePct65up +
            numbUrban +
            pctUrban +
            medIncome +
            pctWWage +
            pctWFarmSelf +
            pctWInvInc +
            pctWSocSec +
            pctWPubAsst +
            pctWRetire +
            medFamInc +
            perCapInc +
            whitePerCap +
            blackPerCap +
            indianPerCap +
            AsianPerCap +
            OtherPerCap +
            HispPerCap +
            NumUnderPov +
            PctPopUnderPov +
            PctLess9thGrade +
            PctNotHSGrad +
            PctBSorMore +
            PctUnemployed +
            PctEmploy +
            PctEmplManu +
            PctEmplProfServ +
            PctOccupManu +
            PctOccupMgmtProf +
            MalePctDivorce +
            MalePctNevMarr +
            FemalePctDiv +
            TotalPctDiv +
            PersPerFam +
            PctFam2Par +
            PctKids2Par +
            PctYoungKids2Par +
            PctTeen2Par +
            PctWorkMomYoungKids +
            PctWorkMom +
            NumIlleg +
            PctIlleg +

```

```

NumImmig +
PctImmigRecent +
PctImmigRec5 +
PctImmigRec8 +
PctImmigRec10 +
PctRecentImmig +
PctRecImmig5 +
PctRecImmig8 +
PctRecImmig10 +
PctSpeakEnglOnly +
PctNotSpeakEnglWell +
PctLargHouseFam +
PctLargHouseOccup +
PersPerOccupHous +
PersPerOwnOccHous +
PersPerRentOccHous +
PctPersOwnOccup +
PctPersDenseHous +
PctHousLess3BR +
MedNumBR +
HousVacant +
PctHousOccup +
PctHousOwnOcc +
PctVacantBoarded +
PctVacMore6Mos +
MedYrHousBuilt +
PctHousNoPhone +
PctWOFullPlumb +
OwnOccLowQuart +
OwnOccMedVal +
OwnOccHiQuart +
RentLowQ +
RentMedian +
RentHighQ +
MedRent +
MedRentPctHousInc +
MedOwnCostPctInc +
MedOwnCostPctIncNoMtg +
NumInShelters +
NumStreet +
PctForeignBorn +
PctBornSameState +
PctSameHouse85 +
PctSameCity85 +
PctSameState85,
data = d, nvmax = 25, method = "forward")

```

```

## Subset selection object
## Call: regsubsets.formula(ViolentCrimesPerPop ~ population + householdsize +
##   racepctblack + racePctWhite + racePctAsian + racePctHisp +
##   agePct12t21 + agePct12t29 + agePct16t24 + agePct65up + numbUrban +
##   pctUrban + medIncome + pctWWage + pctWFarmSelf + pctWInvInc +
##   pctWSocSec + pctWPubAsst + pctWRetire + medFamInc + perCapInc +
##   whitePerCap + blackPerCap + indianPerCap + AsianPerCap +

```

```

##      OtherPerCap + HispPerCap + NumUnderPov + PctPopUnderPov +
##      PctLess9thGrade + PctNotHSGrad + PctBSorMore + PctUnemployed +
##      PctEmploy + PctEmplManu + PctEmplProfServ + PctOccupManu +
##      PctOccupMgmtProf + MalePctDivorce + MalePctNevMarr + FemalePctDiv +
##      TotalPctDiv + PersPerFam + PctFam2Par + PctKids2Par + PctYoungKids2Par +
##      PctTeen2Par + PctWorkMomYoungKids + PctWorkMom + NumIlleg +
##      PctIlleg + NumImmig + PctImmigRecent + PctImmigRec5 + PctImmigRec8 +
##      PctImmigRec10 + PctRecentImmig + PctRecImmig5 + PctRecImmig8 +
##      PctRecImmig10 + PctSpeakEnglOnly + PctNotSpeakEnglWell +
##      PctLargHouseFam + PctLargHouseOccup + PersPerOccupHous +
##      PersPerOwnOccHous + PersPerRentOccHous + PctPersOwnOccup +
##      PctPersDenseHous + PctHousLess3BR + MedNumBR + HousVacant +
##      PctHousOccup + PctHousOwnOcc + PctVacantBoarded + PctVacMore6Mos +
##      MedYrHousBuilt + PctHousNoPhone + PctWOFullPlumb + OwnOccLowQuart +
##      OwnOccMedVal + OwnOccHiQuart + RentLowQ + RentMedian + RentHighQ +
##      MedRent + MedRentPctHousInc + MedOwnCostPctInc + MedOwnCostPctIncNoMtg +
##      NumInShelters + NumStreet + PctForeignBorn + PctBornSameState +
##      PctSameHouse85 + PctSameCity85 + PctSameState85, data = d,
##      nvmax = 25, method = "forward")
## 96 Variables (and intercept)
##
##              Forced in Forced out
## population          FALSE      FALSE
## householdsize        FALSE      FALSE
## racepctblack          FALSE      FALSE
## racePctWhite          FALSE      FALSE
## racePctAsian          FALSE      FALSE
## racePctHisp           FALSE      FALSE
## agePct12t21           FALSE      FALSE
## agePct12t29           FALSE      FALSE
## agePct16t24           FALSE      FALSE
## agePct65up            FALSE      FALSE
## numbUrban             FALSE      FALSE
## pctUrban              FALSE      FALSE
## medIncome             FALSE      FALSE
## pctWWage              FALSE      FALSE
## pctWFarmSelf          FALSE      FALSE
## pctWInvInc            FALSE      FALSE
## pctWSocSec            FALSE      FALSE
## pctWPubAsst           FALSE      FALSE
## pctWRetire            FALSE      FALSE
## medFamInc             FALSE      FALSE
## perCapInc             FALSE      FALSE
## whitePerCap           FALSE      FALSE
## blackPerCap           FALSE      FALSE
## indianPerCap          FALSE      FALSE
## AsianPerCap           FALSE      FALSE
## OtherPerCap           FALSE      FALSE
## HispPerCap            FALSE      FALSE
## NumUnderPov           FALSE      FALSE
## PctPopUnderPov        FALSE      FALSE
## PctLess9thGrade       FALSE      FALSE
## PctNotHSGrad          FALSE      FALSE
## PctBSorMore           FALSE      FALSE
## PctUnemployed         FALSE      FALSE

```

## PctEmploy	FALSE	FALSE
## PctEmplManu	FALSE	FALSE
## PctEmplProfServ	FALSE	FALSE
## PctOccupManu	FALSE	FALSE
## PctOccupMgmtProf	FALSE	FALSE
## MalePctDivorce	FALSE	FALSE
## MalePctNevMarr	FALSE	FALSE
## FemalePctDiv	FALSE	FALSE
## TotalPctDiv	FALSE	FALSE
## PersPerFam	FALSE	FALSE
## PctFam2Par	FALSE	FALSE
## PctKids2Par	FALSE	FALSE
## PctYoungKids2Par	FALSE	FALSE
## PctTeen2Par	FALSE	FALSE
## PctWorkMomYoungKids	FALSE	FALSE
## PctWorkMom	FALSE	FALSE
## NumIlleg	FALSE	FALSE
## PctIlleg	FALSE	FALSE
## NumImmig	FALSE	FALSE
## PctImmigRecent	FALSE	FALSE
## PctImmigRec5	FALSE	FALSE
## PctImmigRec8	FALSE	FALSE
## PctImmigRec10	FALSE	FALSE
## PctRecentImmig	FALSE	FALSE
## PctRecImmig5	FALSE	FALSE
## PctRecImmig8	FALSE	FALSE
## PctRecImmig10	FALSE	FALSE
## PctSpeakEnglOnly	FALSE	FALSE
## PctNotSpeakEnglWell	FALSE	FALSE
## PctLargHouseFam	FALSE	FALSE
## PctLargHouseOccup	FALSE	FALSE
## PersPerOccupHous	FALSE	FALSE
## PersPerOwnOccHous	FALSE	FALSE
## PersPerRentOccHous	FALSE	FALSE
## PctPersOwnOccup	FALSE	FALSE
## PctPersDenseHous	FALSE	FALSE
## PctHousLess3BR	FALSE	FALSE
## MedNumBR	FALSE	FALSE
## HousVacant	FALSE	FALSE
## PctHousOccup	FALSE	FALSE
## PctHousOwnOcc	FALSE	FALSE
## PctVacantBoarded	FALSE	FALSE
## PctVacMore6Mos	FALSE	FALSE
## MedYrHousBuilt	FALSE	FALSE
## PctHousNoPhone	FALSE	FALSE
## PctWOfullPlumb	FALSE	FALSE
## OwnOccLowQuart	FALSE	FALSE
## OwnOccMedVal	FALSE	FALSE
## OwnOccHiQuart	FALSE	FALSE
## RentLowQ	FALSE	FALSE
## RentMedian	FALSE	FALSE
## RentHighQ	FALSE	FALSE
## MedRent	FALSE	FALSE
## MedRentPctHousInc	FALSE	FALSE

```
## MedOwnCostPctInc          FALSE      FALSE
## MedOwnCostPctIncNoMtg     FALSE      FALSE
## NumInShelters             FALSE      FALSE
## NumStreet                 FALSE      FALSE
## PctForeignBorn            FALSE      FALSE
## PctBornSameState          FALSE      FALSE
## PctSameHouse85            FALSE      FALSE
## PctSameCity85             FALSE      FALSE
## PctSameState85            FALSE      FALSE
## 1 subsets of each size up to 25
## Selection Algorithm: forward
```

```
forwardmodelfit <- function(training_data){
  mf <- lm(ViolentCrimesPerPop ~ NumStreet +
    MedOwnCostPctIncNoMtg +
    MedRentPctHousInc +
    MedRent +
    RentLowQ +
    PctHousNoPhone +
    PctVacMore6Mos +
    PctVacantBoarded +
    PctHousOccup +
    PctImmigRec10 +
    PctImmigRec8 +
    PctImmigRec5 +
    PctIlleg +
    PctWorkMom +
    FemalePctDiv +
    MalePctDivorce +
    PctEmplProfServ +
    PctEmploy +
    PctNotHSGrad +
    PctLess9thGrade +
    indianPerCap +
    pctWSocSec +
    pctUrban +
    racePctHisp +
    racePctWhite,
    training_data)
  mf
}
group_F_MSE(forwardmodelfit(d), d)
```

```
## [1] 0.01688268
```

```
backwardmodelfit <- function(training_data){
  mb <- lm(ViolentCrimesPerPop ~ population +
    racePctWhite +
    agePct12t29 +
    numbUrban +
    medIncome +
    PctEmploy +
    pctWSocSec +
    medFamInc +
    whitePerCap +
```

```

PctEmploy +
MalePctDivorce +
MalePctNevMarr +
PctWorkMom +
PctIlleg +
PctImmigRec5 +
PctImmigRec8 +
PersPerOccupHous +
PctHousLess3BR +
PctVacantBoarded +
PctHousOccup +
PctVacMore6Mos +
RentLowQ +
MedRent +
MedOwnCostPctIncNoMtg +
NumStreet,
training_data)
mb
}
group_F_MSE(backwardmodelfit(d), d)

## [1] 0.01684645
group_F_MSE(group_F_fit2.0(group_F_process(d)), group_F_process(d))

## [1] 0.02039485
summary(forwardmodelfit(d))

##
## Call:
## lm(formula = ViolentCrimesPerPop ~ NumStreet + MedOwnCostPctIncNoMtg +
##     MedRentPctHousInc + MedRent + RentLowQ + PctHousNoPhone +
##     PctVacMore6Mos + PctVacantBoarded + PctHousOccup + PctImmigRec10 +
##     PctImmigRec8 + PctImmigRec5 + PctIlleg + PctWorkMom + FemalePctDiv +
##     MalePctDivorce + PctEmplProfServ + PctEmploy + PctNotHSGrad +
##     PctLess9thGrade + indianPerCap + pctWSocSec + pctUrban +
##     racePctHispanic + racePctWhite, data = training_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.50745 -0.07539 -0.01158  0.05254  0.63131
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.007454   0.093093   0.080 0.936206
## NumStreet      0.190690   0.040275   4.735 2.61e-06 ***
## MedOwnCostPctIncNoMtg -0.109093  0.030841  -3.537 0.000428 ***
## MedRentPctHousInc  0.075303   0.038247   1.969 0.049325 *
## MedRent         0.435922   0.090754   4.803 1.87e-06 ***
## RentLowQ       -0.461313   0.090337  -5.107 4.13e-07 ***
## PctHousNoPhone  -0.056889   0.043824  -1.298 0.194628
## PctVacMore6Mos  -0.060606   0.035142  -1.725 0.084998 .
## PctVacantBoarded  0.066917   0.030083   2.224 0.026408 *
## PctHousOccup    -0.105796   0.032763  -3.229 0.001294 **

```

```
## PctImmigRec10      -0.112990    0.076586   -1.475  0.140532
## PctImmigRec8       0.386945    0.110547    3.500  0.000491 ***
## PctImmigRec5      -0.233562    0.074136   -3.150  0.001693 **
## PctIlleg           0.324441    0.046463    6.983  6.24e-12 ***
## PctWorkMom        -0.102374    0.039767   -2.574  0.010227 *
## FemalePctDiv       -0.018782    0.085610   -0.219  0.826409
## MalePctDivorce     0.294348    0.083708    3.516  0.000463 ***
## PctEmplProfServ    0.068993    0.041896    1.647  0.100010
## PctEmploy          0.208276    0.082885    2.513  0.012179 *
## PctNotHSGrad       0.327993    0.110782    2.961  0.003163 **
## PctLess9thGrade    -0.156711    0.084835   -1.847  0.065094 .
## indianPerCap      -0.062586    0.032145   -1.947  0.051900 .
## pctWSocSec         0.169979    0.062461    2.721  0.006647 **
## pctUrban           0.041869    0.013461    3.110  0.001938 **
## racePctHispanic    0.071670    0.034060    2.104  0.035684 *
## racePctWhite       -0.217587    0.040351   -5.392  9.25e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1321 on 774 degrees of freedom
## Multiple R-squared:  0.6937, Adjusted R-squared:  0.6838
## F-statistic: 70.13 on 25 and 774 DF,  p-value: < 2.2e-16
```

```
summary(backwardmodelfit(d))
```

```
##
## Call:
## lm(formula = ViolentCrimesPerPop ~ population + racePctWhite +
##     agePct12t29 + numbUrban + medIncome + PctEmploy + pctWSocSec +
##     medFamInc + whitePerCap + PctEmploy + MalePctDivorce + MalePctNevMarr +
##     PctWorkMom + PctIlleg + PctImmigRec5 + PctImmigRec8 + PersPerOccupHous +
##     PctHousLess3BR + PctVacantBoarded + PctHousOccup + PctVacMore6Mos +
##     RentLowQ + MedRent + MedOwnCostPctIncNoMtg + NumStreet, data = training_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.48416 -0.07536 -0.01504  0.05005  0.64174
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.11404    0.12326   0.925  0.355128
## population     -1.25891    0.35833  -3.513  0.000468 ***
## racePctWhite   -0.24031    0.04495  -5.347  1.18e-07 ***
## agePct12t29    -0.19332    0.09133  -2.117  0.034601 *
## numbUrban       1.29995    0.35654   3.646  0.000284 ***
## medIncome      -0.41544    0.21448  -1.937  0.053113 .
## PctEmploy       0.15422    0.07800   1.977  0.048385 *
## pctWSocSec      0.19251    0.07288   2.642  0.008420 **
## medFamInc       0.56811    0.20508   2.770  0.005737 **
## whitePerCap    -0.24261    0.09689  -2.504  0.012488 *
## MalePctDivorce  0.29898    0.04932   6.062  2.10e-09 ***
## MalePctNevMarr  0.12344    0.06899   1.789  0.073955 .
## PctWorkMom     -0.10614    0.04140  -2.564  0.010548 *
## PctIlleg        0.34401    0.04835   7.114  2.56e-12 ***
## PctImmigRec5   -0.22614    0.07289  -3.103  0.001988 **
```




```
## PctImmigRec8      0.28487    0.07858    3.625 0.000307 ***
## PersPerOccupHous  0.16902    0.06770    2.497 0.012740 *
## PctHousLess3BR    0.08395    0.05506    1.525 0.127770
## PctVacantBoarded  0.07285    0.03106    2.345 0.019279 *
## PctHousOccup      -0.12161    0.03311   -3.673 0.000256 ***
## PctVacMore6Mos    -0.07693    0.03448   -2.231 0.025964 *
## RentLowQ          -0.47068    0.09126   -5.158 3.18e-07 ***
## MedRent           0.46991    0.10723    4.382 1.34e-05 ***
## MedOwnCostPctIncNoMtg -0.10267    0.03000   -3.422 0.000653 ***
## NumStreet         0.17518    0.04918    3.562 0.000391 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1319 on 775 degrees of freedom
## Multiple R-squared:  0.6944, Adjusted R-squared:  0.6849
## F-statistic: 73.37 on 24 and 775 DF,  p-value: < 2.2e-16
```

#### *#Data Exploration*

```
d <- read.csv("http://andrewpbray.github.io/data/crime-train.csv")
```


#### *#maybe*

```
ggplot(data = d, mapping = aes(x = householdsize, y = ViolentCrimesPerPop)) +
  geom_point(alpha = .5)
```




*#ok*

```
ggplot(data = d, mapping = aes(x = PctUnemployed, y = ViolentCrimesPerPop)) +  
  geom_point(alpha = .5)
```




*#ok*


```
ggplot(data = d, mapping = aes(x = medIncome, y = ViolentCrimesPerPop)) +  
  geom_point(alpha = .5)
```



```
#ok  
ggplot(data = d, mapping = aes(x = PctPopUnderPov, y = ViolentCrimesPerPop)) +  
  geom_point(alpha = .5)
```




```
#ok  
ggplot(data = d, mapping = aes(x = PctNotHSGrad, y = ViolentCrimesPerPop)) +  
  geom_point(alpha = .5)
```




```
#ok  
ggplot(data = d, mapping = aes(x = PctKids2Par, y = ViolentCrimesPerPop)) +  
  geom_point(alpha = .5) +  
  geom_smooth()
```


```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



```
ggplot(data = d, mapping = aes(x = LemasPctOfficDrugUn, y = ViolentCrimesPerPop)) +  
  geom_point(alpha = .5)
```



```
ggplot(data = d, mapping = aes(x = NumStreet, y = ViolentCrimesPerPop)) +  
  geom_point(alpha = .5)
```



```
cor(d$PctKids2Par, d$PctNotHSGrad) #-.697
```

```
## [1] -0.6974628
```

```
cor(d$PctKids2Par, d$PctPopUnderPov) #-.771
```

```
## [1] -0.7713265
```

```
cor(d$PctKids2Par, d$medIncome) #.712
```

```
## [1] 0.7120151
```

```
cor(d$PctKids2Par, d$PctUnemployed) #-.698
```

```
## [1] -0.6979051
```

```
cor(d$ViolentCrimesPerPop, d$PctPopUnderPov) #.538
```

```
## [1] 0.5377977
```

```
cor(d$ViolentCrimesPerPop, d$medIncome) #-.414
```

```
## [1] -0.4138494
```

```
cor(d$ViolentCrimesPerPop, d$PctFam2Par) #-.705
```

```
## [1] -0.7052806
```

```
cor(d$ViolentCrimesPerPop, d$householdsize) #-0.386
```

```
## [1] -0.03861389
```



```

cor(d$ViolentCrimesPerPop, d$PctUnemployed) #.537

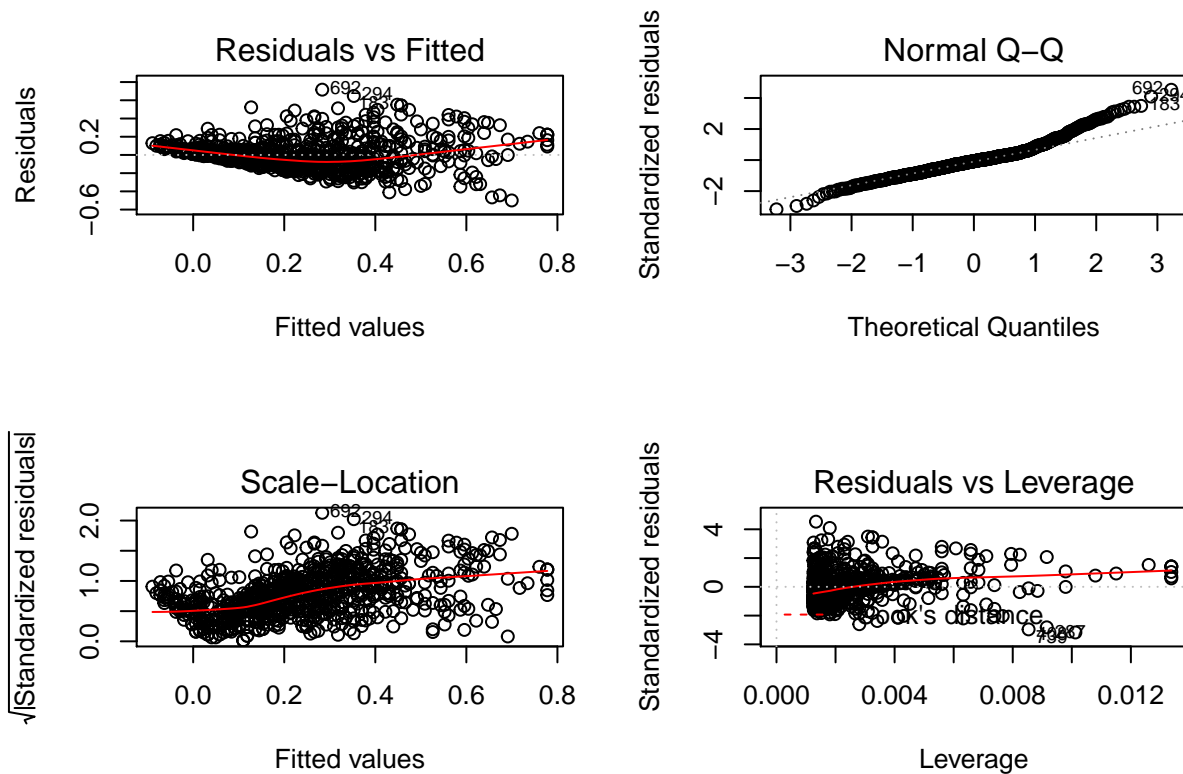
## [1] 0.5373199
cor(d$ViolentCrimesPerPop, d$PctNotHSGrad) #.5157

## [1] 0.5156772
cor(d$PctUnemployed, d$PctPopUnderPov) #.7839

## [1] 0.7839299
#r2
modelkd <- lm(ViolentCrimesPerPop ~ PctKids2Par + PctPopUnderPov, d)
summary(modelkd)

##
## Call:
## lm(formula = ViolentCrimesPerPop ~ PctKids2Par + PctPopUnderPov,
##     data = d)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.47643 -0.10037 -0.01265  0.06511  0.73254
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.84855    0.03776   22.472  <2e-16 ***
## PctKids2Par   -0.93959    0.04383  -21.436  <2e-16 ***
## PctPopUnderPov -0.08756    0.04086   -2.143   0.0324 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1579 on 797 degrees of freedom
## Multiple R-squared:  0.5492, Adjusted R-squared:  0.548
## F-statistic: 485.4 on 2 and 797 DF, p-value: < 2.2e-16
#r2 = .5656
model1 <- lm(ViolentCrimesPerPop ~ PctKids2Par, data = d)
par(mfrow = c(2,2))
plot(model1)

```




```
summary(model1)
```


```
##
## Call:
## lm(formula = ViolentCrimesPerPop ~ PctKids2Par, data = d)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.49968 -0.09737 -0.01168  0.06528  0.71654
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.77773     0.01830   42.49  <2e-16 ***
## PctKids2Par -0.86713     0.02796  -31.01  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1583 on 798 degrees of freedom
## Multiple R-squared:  0.5466, Adjusted R-squared:  0.546
## F-statistic: 961.9 on 1 and 798 DF, p-value: < 2.2e-16
```

```
ggplot(data = d, mapping = aes(x = racePctWhite, y = ViolentCrimesPerPop)) +
  geom_point(alpha = .5) + geom_smooth()
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```




```
ggplot(data = d, mapping = aes(x = pctUrban, y = ViolentCrimesPerPop)) +  
  geom_point(alpha = .5)
```




```
ggplot(data = d, mapping = aes(x = MalePctDivorce, y = ViolentCrimesPerPop)) +  
  geom_point(alpha = .5) +  
  geom_smooth()
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```




```
ggplot(data = d, mapping = aes(x = PctKids2Par, y = ViolentCrimesPerPop)) +  
  geom_point(alpha = .5) +  
  geom_smooth()
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```




```
ggplot(data = d, mapping = aes(x = PctIlleg, y = ViolentCrimesPerPop)) +  
  geom_point(alpha = .5) +  
  geom_smooth()
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



```
ggplot(data = d, mapping = aes(x = PctPersDenseHous, y = ViolentCrimesPerPop)) +  
  geom_point(alpha = .5) +  
  geom_smooth()
```

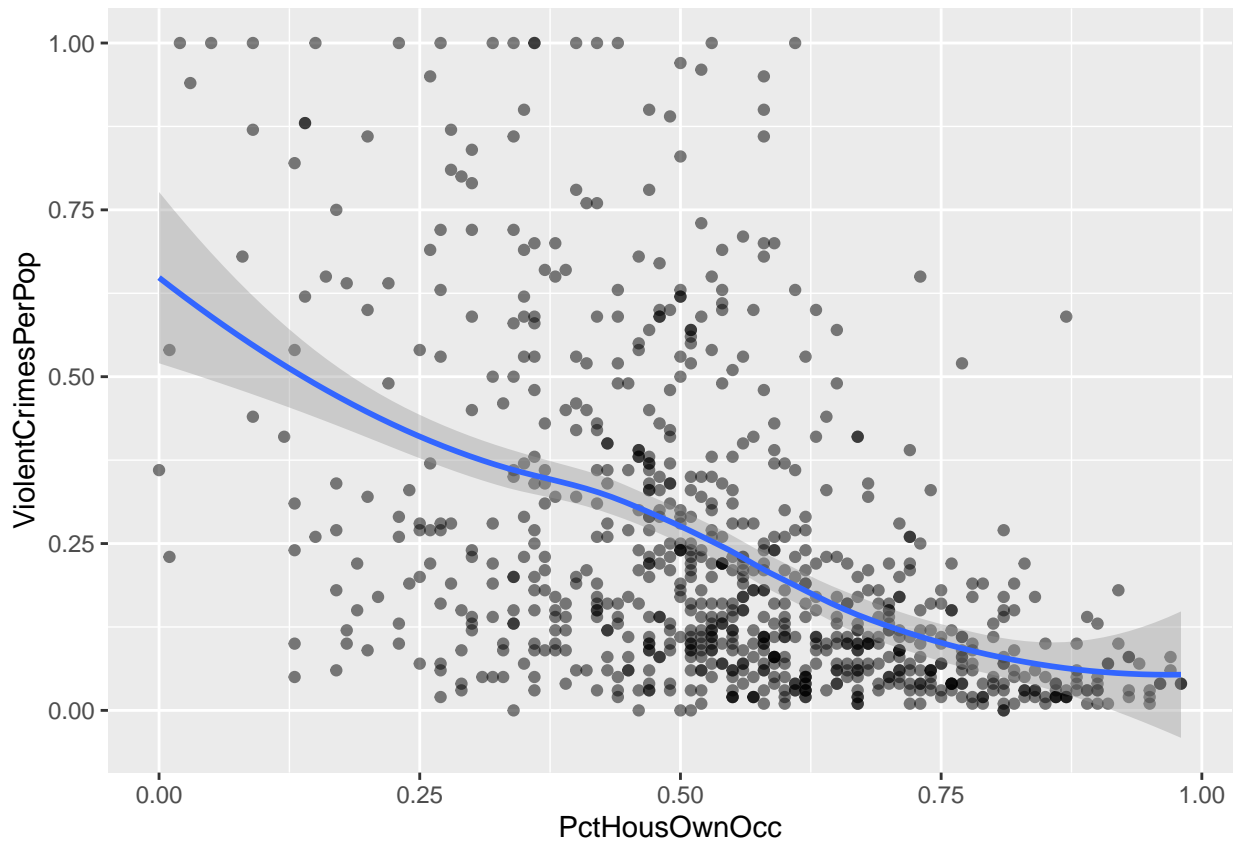
```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



```
ggplot(data = d, mapping = aes(x = PctHousOwnOcc, y = ViolentCrimesPerPop)) +  
  geom_point(alpha = .5) +  
  geom_smooth()
```


```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```





```
ggplot(data = d, mapping = aes(x = NumStreet, y = ViolentCrimesPerPop)) +
  geom_point(alpha = .5) +
  geom_smooth()
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : at -0.005
## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : radius 2.5e-05
## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : all data on boundary of neighborhood. make span bigger
## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : pseudoinverse used at -0.005
## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : neighborhood radius 0.005
## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : reciprocal condition number 1
## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : zero-width neighborhood. make span bigger
## Warning: Computation failed in `stat_smooth()`:
## NA/NaN/Inf in foreign function call (arg 5)
```



```
#I(MalePctDivorce^(1/3))+I(PctKids2Par^(1/3)) + PctIlleg+(PctPersDenseHous^2) + RentLowQ+MedRent+NumStr
```

```
library(tidyverse)
library(leaps)
```

```
d <- read.csv("http://andrewpbray.github.io/data/crime-train.csv")
```

```
# Data wrangling
```

```
group_F_process <- function(training_data) {
```

```
  dw<-as_tibble(training_data) %>%
```

```
    mutate(
```

```
      MalePctDivorceCub = I(MalePctDivorce^(1/3)),
```

```
      PctKids2ParCub = I(PctKids2Par^(1/3)),
```

```
      PctPersDenseHousSq = PctPersDenseHous^2
```

```
    )
```

```
  dw<-select (d,-c(state,county,community,communityname,LemasSwornFT,LemasSwFTPerPop,LemasSwFTFieldOps,
```

```
  vars<-c()
```

```
  for (i in 1:(length(dw)-1)) {
```

```
    vars<- c(vars, names(dw)[i])
```

```
  }
```

```
  sqrd<-data.frame(lapply(vars, function(x){dw[,x]^(1/2)}))
```

```
  cubc<-data.frame(lapply(vars, function(x){dw[,x]^(1/3)}))
```

```
  names(sqrd)<-paste0(vars, "Sq")
```

```
  names(cubc)<-paste0(vars, "Cub")
```

```
  dw<-cbind(dw, sqrd,cubc)
```

```
  return(dw)
```

```
}
```

```

# Manually fits model
group_F_fit <- function(training_data) {
  m1 <- lm(ViolentCrimesPerPop~
           MalePctDivorceCub+
           PctKids2ParCub+
           PctIlleg+
           PctPersDenseHousSq+
           RentLowQ+
           MedRent,
           training_data)

  m1
}

# Gets MSE
group_F_MSE <- function(model, data) {
  mean((data$ViolentCrimesPerPop - predict.lm(model, data)) ^ 2)
}

# Automatically fits model
group_F_automated_fit <- function(data, method) {
  leaps<-regsubsets(ViolentCrimesPerPop~.,
                    data = data,
                    nvmax = 25,
                    method = method)

  best<-summary(leaps)$which[which.max(summary(leaps)$adjr2),]
  variables <- c()
  for (i in 2:length(best)) {
    if (best[i] == TRUE) {
      variables <- c(variables, names(best)[i])
    }
  }
  vars<- paste(variables, collapse = "+")
  formula<- paste("lm(ViolentCrimesPerPop ~ ",vars," , data = dw)") # not able to get dataframe name as
  m1<-eval(parse(text=formula))
  return(m1)
}

dw <- group_F_process(d)
bestF <- group_F_automated_fit(dw, "forward")

## Warning in leaps.setup(x, y, wt = wt, nbest = nbest, nvmax = nvmax,
## force.in = force.in, : 1 linear dependencies found
## Reordering variables and trying again:
bestB <- group_F_automated_fit(dw, "backward")

## Warning in leaps.setup(x, y, wt = wt, nbest = nbest, nvmax = nvmax,
## force.in = force.in, : 1 linear dependencies found
## Reordering variables and trying again:
MSE_F <- group_F_MSE(bestF, dw)
MSE_B <- group_F_MSE(bestB, dw)

```